

# Databases

Lecture 7

Relational Algebra

- query languages in the relational model
  - relational algebra and calculus - formal query languages with a significant influence on SQL
    - relational algebra
      - queries are specified in an operational manner
    - relational calculus
      - queries describe the desired answer, without specifying how it will be computed (declarative)
  - not expected to be Turing complete
  - not intended for complex calculations
  - provide efficient access to large datasets
  - allow optimizations

- relational algebra
  - used by DBMSs to represent query execution plans
  - a relational algebra query:
    - is built using a collection of operators
    - describes a step-by-step procedure for computing the result set
    - is evaluated on the input relations' instances
    - produces an instance of the output relation
  - every operation returns a relation, so operators can be composed; the algebra is closed
  - the result of an algebra expression is a relation, and a relation is a set of tuples
- relational algebra on bags (multisets) - duplicates are not eliminated

## Conditions

- conditions that can be used in several algebraic operators
- similar to the SELECT filter conditions

### 1. *attribute\_name relational\_operator value*

- *value* - attribute name, expression

### 2. *attribute\_name IS [NOT] IN single\_column\_relation*

- a relation with one column can be considered a set
- the condition tests whether a value belongs to a set

### 3. *relation {IS [NOT] IN | = | <>} relation*

- the relations in the condition must be union-compatible

## Conditions

4. *(condition)*

*NOT condition*

*condition<sub>1</sub> AND condition<sub>2</sub>*

*condition<sub>1</sub> OR condition<sub>2</sub>,*

where *condition*, *condition<sub>1</sub>*, *condition<sub>2</sub>* are conditions of type 1-4.

## Operators in the Algebra

- equivalent SELECT statements can be specified for the relational algebra expressions
- *selection*
  - notation:  $\sigma_C(R)$
  - resulting relation:
    - schema:  $R$ 's schema
    - tuples: records in  $R$  that satisfy condition  $C$
  - equivalent SELECT statement

```
SELECT *  
FROM R  
WHERE C
```

- *projection*
  - notation:  $\pi_{\alpha}(R)$
  - resulting relation:
    - schema: attributes in  $\alpha$
    - tuples: every record in  $R$  is projected on  $\alpha$
  - $\alpha$  can be extended to a set of expressions, specifying the columns of the relation being computed
  - equivalent SELECT statement

```
SELECT DISTINCT  $\alpha$ 
FROM R
```

```
SELECT  $\alpha$ 
FROM R                -- algebra on bags
```

- *cross-product*
  - notation:  $R_1 \times R_2$
  - resulting relation:
    - schema: the attributes of  $R_1$  followed by the attributes of  $R_2$
    - tuples: every tuple  $r_1$  in  $R_1$  is concatenated with every tuple  $r_2$  in  $R_2$
- equivalent SELECT statement

```
SELECT *  
FROM R1 CROSS JOIN R2
```



- *union, set-difference, intersection*
  - notation:  $R_1 \cup R_2$ ,  $R_1 - R_2$ ,  $R_1 \cap R_2$
  - $R_1$  and  $R_2$  must be union-compatible:
    - same number of columns
    - corresponding columns, taken in order from left to right, have the same domains
  - equivalent SELECT statements

SELECT *	SELECT *	SELECT *
FROM R1	FROM R1	FROM R1
UNION	EXCEPT	INTERSECT
SELECT *	SELECT *	SELECT *
FROM R2	FROM R2	FROM R2

-- algebra on bags: SELECT statements that don't eliminate duplicates (e.g., UNION ALL)

- join operators
  - *condition join* (or *theta join*)
    - notation:  $R_1 \otimes_{\Theta} R_2$
    - result: the records in the cross-product of  $R_1$  and  $R_2$  that satisfy a certain condition
  - definition  $\Rightarrow R_1 \otimes_{\Theta} R_2 = \sigma_{\Theta}(R_1 \times R_2)$
  - equivalent SELECT statement
 

```
SELECT *
FROM R1 INNER JOIN R2 ON  $\Theta$ 
```

- join operators
  - *natural join*
    - notation:  $R_1 * R_2$
    - resulting relation:
      - schema: the union of the attributes of the two relations (attributes with the same name in  $R_1$  and  $R_2$  appear once in the result)
      - tuples: obtained from tuples  $\langle r_1, r_2 \rangle$ , where  $r_1$  in  $R_1$ ,  $r_2$  in  $R_2$ , and  $r_1$  and  $r_2$  agree on the common attributes of  $R_1$  and  $R_2$
  - let  $R_1[\alpha]$ ,  $R_2[\beta]$ ,  $\alpha \cap \beta = \{A_1, A_2, \dots, A_m\}$ ; then:
 
$$R_1 * R_2 = \pi_{\alpha \cup \beta} (R_1 \otimes_{R_1.A_1=R_2.A_1 \text{ AND } \dots \text{ AND } R_1.A_m=R_2.A_m} R_2)$$
  - equivalent SELECT statement
 

```
SELECT *
FROM R1 NATURAL JOIN R2
```

- join operators
  - *left outer join*
    - notation (in these notes):  $R_1 \bowtie_C R_2$
    - resulting relation:
      - schema: the attributes of  $R_1$  followed by the attributes of  $R_2$
      - tuples: tuples from the condition join  $R_1 \bowtie_C R_2$  + the tuples in  $R_1$  that were not used in  $R_1 \bowtie_C R_2$  combined with the *null* value for the attributes of  $R_2$
  - equivalent SELECT statement

```
SELECT *  
FROM R1 LEFT OUTER JOIN R2 ON C
```

- join operators
  - *right outer join*
    - notation:  $R_1 \bowtie_C R_2$
    - resulting relation:
      - schema: the attributes of  $R_1$  followed by the attributes of  $R_2$
      - tuples: tuples from the condition join  $R_1 \Join_C R_2$  + the tuples in  $R_2$  that were not used in  $R_1 \Join_C R_2$  combined with the *null* value for the attributes of  $R_1$
  - equivalent SELECT statement

```
SELECT *  
FROM R1 RIGHT OUTER JOIN R2 ON C
```

- join operators
  - *full outer join*
    - notation:  $R_1 \bowtie_c R_2$
    - resulting relation:
      - schema: the attributes of  $R_1$  followed by the attributes of  $R_2$
      - tuples:
        - tuples from the condition join  $R_1 \otimes_c R_2$  +
        - the tuples in  $R_1$  that were not used in  $R_1 \otimes_c R_2$  combined with the *null* value for the attributes of  $R_2$  +
        - the tuples in  $R_2$  that were not used in  $R_1 \otimes_c R_2$  combined with the *null* value for the attributes of  $R_1$
  - equivalent SELECT statement

```
SELECT *
FROM R1 FULL OUTER JOIN R2 ON C
```

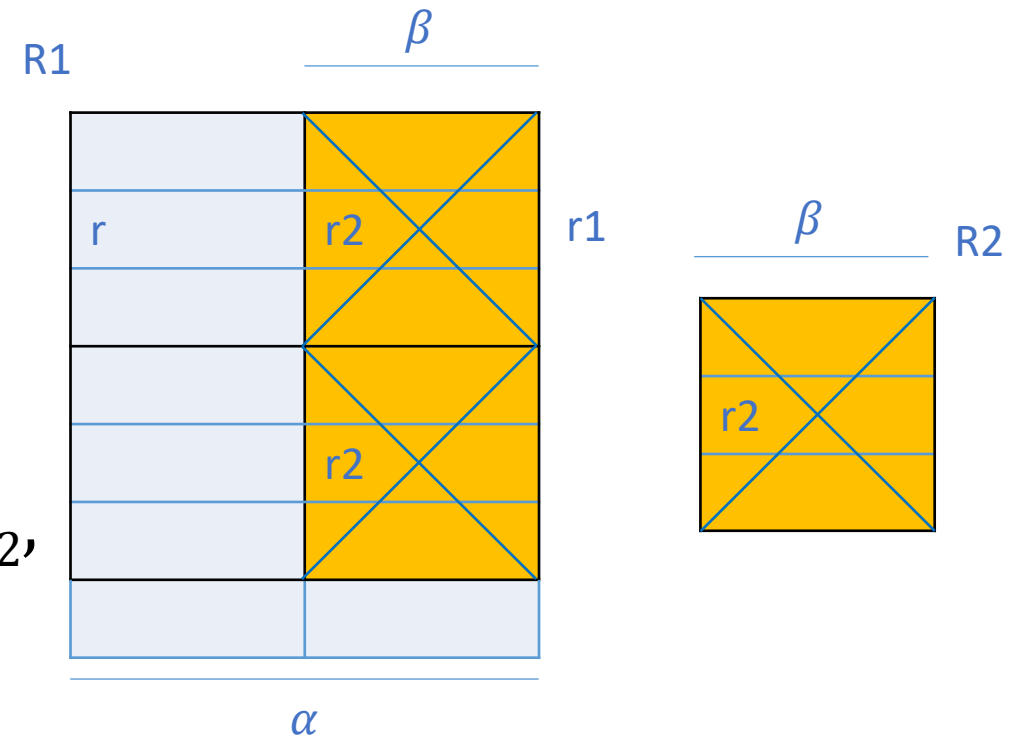
- join operators
  - *left semi join*
    - notation:  $R_1 \triangleright R_2$
    - resulting relation:
      - schema:  $R_1$ 's schema
      - tuples: the tuples in  $R_1$  that are used in the natural join  $R_1 * R_2$

- join operators
  - *right semi join*
    - notation:  $R_1 \Join R_2$
    - resulting relation:
      - schema:  $R_2$ 's schema
      - tuples: the tuples in  $R_2$  that are used in the natural join  $R_1 * R_2$



- *division*

- notation:  $R_1 \div R_2$
- $R_1[\alpha], R_2[\beta], \beta \subset \alpha$
- resulting relation:
  - schema:  $\alpha - \beta$
  - tuples: a record  $r \in R_1 \div R_2$  iff  $\forall r_2 \in R_2, \exists r_1 \in R_1$  such that:
    - $\pi_{\alpha-\beta}(r_1) = r$
    - $\pi_{\beta}(r_1) = r_2$
  - i.e., a record  $r$  belongs to the result if in  $R_1$   $r$  is concatenated with every record in  $R_2$



- see lecture examples (at the board) with algebra queries:
  - selection
  - projection
  - division
  - selection, projection
  - natural join, selection, projection
  - set-difference, natural join, selection, projection
  - different algebra expressions producing the same result (optimization - reducing the size of intermediate relations)

## An Independent Subset of Operators

- independent set of operators M:
  - eliminating any operator  $op$  from M: there will be a relation that can be obtained using M's operators, but cannot be obtained with the operators in  $M - \{op\}$
- for the previously described query language, with operators:  
 $\{\sigma, \pi, \times, \cup, -, \cap, \otimes, *, \ltimes, \rtimes, \bowtie, \triangleright, \triangleleft, \div\}$   
an independent set of operators is  $\{\sigma, \pi, \times, \cup, -\}$
- the other operators are obtained as follows (some expressions have already been introduced):
  - $R_1 \cap R_2 = R_1 - (R_1 - R_2)$
  - $R_1 \otimes_C R_2 = \sigma_C(R_1 \times R_2)$

- the other operators are obtained as follows (some expressions have already been introduced):

- $R_1[\alpha], R_2[\beta], \alpha \cap \beta = \{A_1, A_2, \dots, A_m\}$ , then:

$$R_1 * R_2 = \pi_{\alpha \cup \beta}(R_1 \otimes_{R_1.A_1=R_2.A_1 \text{ AND } \dots \text{ AND } R_1.A_m=R_2.A_m} R_2)$$

- $R_1[\alpha], R_2[\beta], R_3[\beta] = \{(null, \dots, null)\}, R_4[\alpha] = \{(null, \dots, null)\}$

$$R_1 \bowtie_C R_2 = (R_1 \otimes_C R_2) \cup (R_1 - \pi_\alpha(R_1 \otimes_C R_2)) \times R_3$$

$$R_1 \bowtie_C R_2 = (R_1 \otimes_C R_2) \cup R_4 \times (R_2 - \pi_\beta(R_1 \otimes_C R_2))$$

$$R_1 \bowtie_C R_2 = (R_1 \bowtie_C R_2) \cup (R_1 \bowtie_C R_2)$$

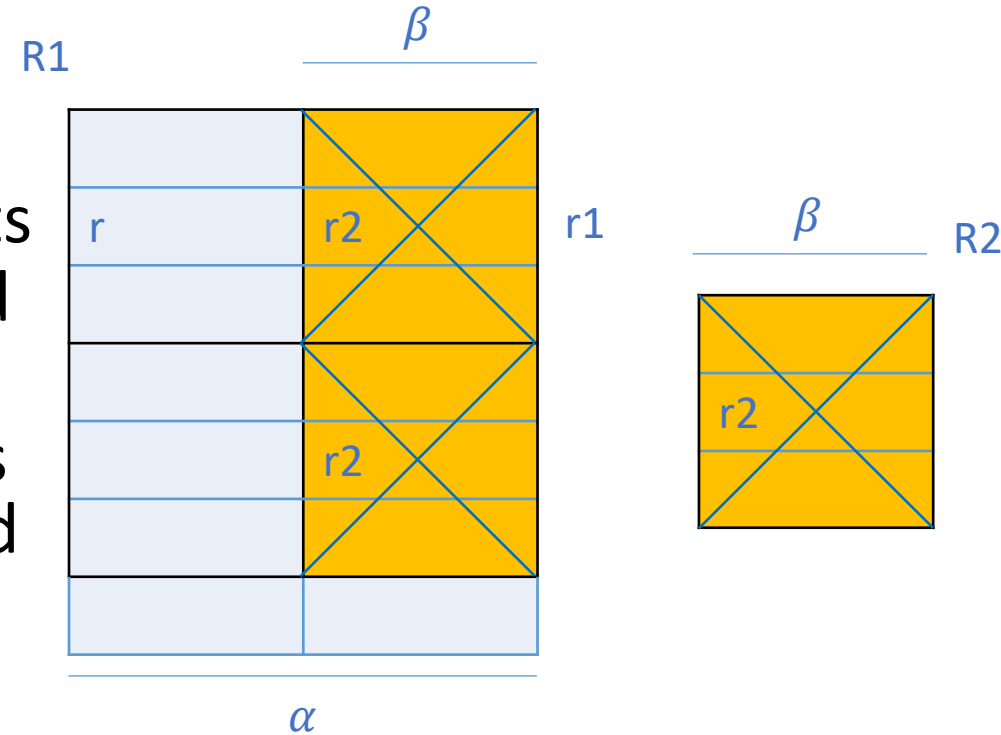
- $R_1[\alpha], R_2[\beta]$

$$R_1 \triangleright R_2 = \pi_\alpha(R_1 * R_2)$$

$$R_1 \triangleleft R_2 = \pi_\beta(R_1 * R_2)$$

- the other operators are obtained as follows (some expressions have already been introduced):
  - if  $R_1[\alpha]$ ,  $R_2[\beta]$ ,  $\beta \subset \alpha$ , then  $r \in R_1 \div R_2$  iff  $\forall r_2 \in R_2, \exists r_1 \in R_1$  such that:  $\pi_{\alpha-\beta}(r_1) = r$  and  $\pi_\beta(r_1) = r_2$   
 $\Rightarrow r$  is in  $\pi_{\alpha-\beta}(R_1)$ , but not all the elements in  $\pi_{\alpha-\beta}(R_1)$  are in the result
  - $(\pi_{\alpha-\beta}(R_1)) \times R_2$  contains all the elements with one part in  $\pi_{\alpha-\beta}(R_1)$  and the second part in  $R_2$
  - to obtain values that are disqualified,  $R_1$  is subtracted from the obtained relation, and the result is projected on  $\alpha - \beta$
  - the final expression:

$$R_1 \div R_2 = \pi_{\alpha-\beta}(R_1) - \pi_{\alpha-\beta}((\pi_{\alpha-\beta}(R_1)) \times R_2 - R_1)$$



# References

- [Ta13] ȚÂMBULEA, L., Curs Baze de date, Facultatea de Matematică și Informatică, UBB, 2013-2014
- [Ra00] RAMAKRISHNAN, R., GEHRKE, J., Database Management Systems (3rd Edition), McGraw-Hill, 2002
- [Da03] DATE, C.J., An Introduction to Database Systems (8<sup>th</sup> Edition), Addison-Wesley, 2003
- [Ga08] GARCIA-MOLINA, H., ULLMAN, J., WIDOM, J., Database Systems: The Complete Book (2nd Edition), Pearson Education, 2009
- [Ra07] RAMAKRISHNAN, R., GEHRKE, J., Database Management Systems, McGraw-Hill, 2007,  
<http://pages.cs.wisc.edu/~dbbook/openAccess/thirdEdition/slides/slides3ed.html>
- [Ul11] ULLMAN, J., WIDOM, J., A First Course in Database Systems,  
<http://infolab.stanford.edu/~ullman/fcdb.html>